

## **Loan Approval Classification Dataset:**

The Loan Approval Classification Dataset contains 45,000 records and 14 variables, designed to model and predict whether a loan application will be approved or rejected. The dataset includes both numerical variables (person's age, income, employment length, loan amount, interest rate, credit score) and categorical variables are home ownership status, loan purpose, loan grade, credit default history. The target variable is loan\_status, a binary indicator where 1 = approved and 0 = rejected.

The sample dataset is derived from the Loan Approval Classification dataset and contains 201 records with a mix of numerical and categorical variables related to loan applications. Unlike the full dataset, this sample includes 17 missing values distributed across five variables: person\_age (4), person\_gender (4), person\_income (4), person\_loan\_status (3), and person\_education (2). Our goal is to apply data preprocessing techniques, handling missing values and exploratory data analysis.

## **Code Description:**

```
install.packages("dplyr")
library(dplyr)
library(ggplot2)
```

Installed and loaded the “dplyr” and “ggplot2” packages for data manipulation and visualization.

### **Code:**

```
mydata<-read.csv("C:/R/Midterm_Dataset_Section(A).csv",header=TRUE,sep="")
print(mydata)

  person_age person_gender person_education person_income
1         21        female       Master      71948
2         21        female   High School     12282
3         25        female   High School     12438
4         23        female      Bachelor     79753
5         24         male       Master      66135
6        NA        female   High School     12951
7         22        female      Bachelor      NA
8         24             High School     95550
9         22        female   High School    100684
10        21        female   High School     12739
```

This code reads mydata.csv dataset into the R studio as mydata and prints all rows and columns.

### **Code:**

```
colSums(is.na(mydata) | mydata == "" )
```

```
> colSums(is.na(mydata) | mydata == "" )
   person_age          person_gender          person_education        person_income
                  4                      4                      2                      4
   person_emp_exp      person_home_ownership      person_loan_amnt        loan_intent
                  0                      0                      0                      0
   loan_int_rate      loan_percent_income      cb_person_cred_hist_length credit_score
                  0                      0                      0                      0
previous_loan_defaults_on_file    loan_status
                  0                      3
```

This code counts the number of missing values in each column for both NA and empty strings ("") as missing as missing value.

### Code:


```
colSums(is.na(mydata) | mydata == "" )

summary(mydata)

missing_data <- data.frame(
  variable = names(mydata),
  missing_count = sapply(mydata, function(x) sum(is.na(x) | x == ""))
)
```

**Counts missing values in each column, considering both NA and empty strings ("")**

```
ggplot(missing_data, aes(x = reorder(variable, -missing_count), y = missing_count)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  coord_flip() + labs(title = "Missing (NA or Empty String) Values per Variable", x = "Variables",
y = "Number of Missing Values")
```



Visualize the number of missing values in a bar plot.

### Code:

```
duplicated(mydata)
```

```
> duplicated(mydata)
[1] FALSE FALSE
[21] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[41] FALSE FALSE
[61] FALSE FALSE
[81] FALSE FALSE
[101] FALSE FALSE
[121] FALSE FALSE
[141] FALSE FALSE
[161] FALSE FALSE
[181] FALSE FALSE
[201] FALSE
> |
```

If there are any duplicate row it will show TRUE, else it will show FALSE

### Code:

```
mean_age <- mean(mydata$person_age, na.rm = TRUE)
```

```
mean_age <- as.integer(round(mean_age))
```

```
mydata$person_age[is.na(mydata$person_age)] <- mean_age
```

- ➔ This code calculates mean value with mean() function. Then fills missing person\_age values with the rounded mean age.

```
mode_gender <- names(sort(table(mydata$person_gender), decreasing = TRUE))[1]
```

```
mydata$person_gender[mydata$person_gender == ""] <- mode_gender
```

- ➔ This code calculates the most frequent gender (mode) and replaces all empty strings in the person\_gender column with this mode.

```
mode_education <- names(sort(table(mydata$person_education), decreasing = TRUE))[1]
```

```
mydata$person_education[mydata$person_education == ""] <- mode_education
```

- ➔ This code finds the most frequent value (mode) in person\_education and replaces all empty strings in that column with the mode.

```
mean_income <- mean(mydata$person_income, na.rm = TRUE)
```

```
mean_income <- as.integer(round(mean_income))
```

```
mydata$person_income[is.na(mydata$person_income)] <- mean_income
```

- ➔ This code computes the mean of person\_income rounds it to the nearest integer, and replaces all missing person\_income with this mean.

```

unique(mydata$loan_status)

mode_loan_status <- names(sort(table(mydata$loan_status),decreasing = TRUE))[1]

mydata$loan_status[is.na(mydata$loan_status)] <- mode_loan_status

```

→ This code finds the mode in loan\_status and replaces with the mode.

After handling missing value run,

```
colSums(is.na(mydata) | mydata == "")
```

```


> colSums(is.na(mydata) | mydata == "") 
    person_age           person_gender          person_education      person_income
                  0                   4                   2                   4
    person_emp_exp       person_home_ownership   loan_amnt
                  0                   0                   0
    loan_int_rate        loan_percent_income   cb_person_cred_hist_length
                  0                   0                   0
previous_loan_defaults_on_file      loan_status
                  0                   3
> |

```

No missing value in our dataset.

## Code: Outliers

```
boxplot(mydata$person_age)
```



→ 4 outliers in person\_age

```
quantile(mydata$person_age)
```

```

> quantile(mydata$person_age)
  0%   25%   50%   75% 100%
 -25    22    23    25   350
> |

```

→ Q1=22, Q3=25

```
x <- mydata$person_age
```

```
x
```

```

iqr <- IQR(x)

iqr

lower_bound_x <- 22 - 1.5 * iqr

upper_bound_x <- 25 + 1.5 * iqr

x[x < lower_bound_x | x > upper_bound_x] <- mean(x)


x <- as.integer(round(x))

mydata$person_age <- x

```


→ This code calculates the IQR of person\_age, determines outlier bounds, replaces values outside these bounds with the mean

```
boxplot(mydata$person_age)
```



→ Outliers removed

```
boxplot(mydata$person_emp_exp)
```



→ person\_emp\_exp has three outlier

```
boxplot(mydata$person_emp_exp)

quantile(mydata$person_emp_exp)

exp <- mydata$person_emp_exp
```

```

exp

iqr_exp <- IQR(exp)

iqr_exp

lower_bound_exp <- 0 - 1.5 * iqr_exp

upper_bound_exp <- 3 + 1.5 * iqr_exp

exp[exp < lower_bound_exp | exp > upper_bound_exp] <- mean(exp)


exp <- as.integer(round(exp))

mydata$person_emp_exp <- exp

```


→ This code calculates the IQR of person\_emp\_exp, determines outlier bounds, replaces values outside these bounds with the mean

```
boxplot(mydata$person_emp_exp)
```



→ Outliers removed

```
boxplot(mydata$person_income)
```



-> person\_income has outlier


```

quantile(mydata$person_income)
i <- mydata$person_income
iqr_i <- IQR(i)
iqr_i
lower_bound_i <- 60747 - 1.5 * iqr_i
upper_bound_i <- 149875 + 1.5 * iqr_i
mean(i)
i[i < lower_bound_i | i > upper_bound_i] <- mean(i)
i <- as.integer(round(i))
mydata$person_income<-i

```

➔ This code calculates the IQR of person\_income, determines outlier bounds, replaces values outside these bounds with the mean

```
boxplot($person_income)
```



➔ Outlier removed

### Code: Noisy Value

```

unique(mydata$person_gender)
> unique(mydata$person_gender)
[1] "female"   "male"     ""          "malee"     "feemale"

```

➔ Here “malee” and “feemale” are noisy value

```
mydata <- mydata %>%
  mutate(gender = case_when(
    tolower(person_gender) %in% c("male", "malee") ~ "male",
    tolower(person_gender) %in% c("female", "feemale") ~ "female"
  ))
```

```
mydata$person_gender<-mydata$gender
```

```
mydata$gender <- NULL
```

➔ This code creates a new Gender column, convert and store misspelled entries into valid format. After that it replaced the values of person\_gende with Gender column. Then replaced the value of gender column with Null and Gender column will be removed.

```
unique(mydata$person_gender)
```

```
> unique(mydata$person_gender)
[1] "female" "male"                                     -> Noisy value replaced
```

```
unique(mydata$person_home_ownership)
```

```
> unique(mydata$person_home_ownership)
[1] "RENT"      "OWN"       "MORTGAGE"   "RENTT"     "OOWN"      "OTHER"
> |
```

➔ RENTT , OOWN are noisy value

```
unique(mydata$person_home_ownership)
```

```
mydata <- mydata %>%
```

```
  mutate(home = case_when(
    toupper(person_home_ownership) %in% c("RENT", "RENTT") ~ "RENT",
```

```

toupper(person_home_ownership) %in% c("OWN", "OOWN") ~ "OWN",
toupper(person_home_ownership) %in% c("MORTGAGE") ~ "MORTGAGE",
toupper(person_home_ownership) %in% c("OTHER") ~ "OTHER" ))
mydata$person_home_ownership<-mydata$home
mydata$home <- NULL

```

→ This code creates a new Homecollum, convert and store misspelled entries into valid format. After that it replaced the values of person\_home\_ownership with Home collum. Then replaced the value of Home collum with Null and Gender collum will be removed.



```

unique(mydata$person_home_ownership)
- - -
> unique(mydata$person_home_ownership)
[1] "RENT"      "OWN"       "MORTGAGE"   "OTHER"     -> Noisy value removed

```

## Categorical to Numeric:

```

mydata$person_gender <- factor(mydata$person_gender,
                                levels = c("female", "male"), labels = c(1,2))
> unique(mydata$person_gender)
[1] 1 2
> unique(mydata$person_gender)
[1] "female" "male"          Levels: 1 2

```

→ Converted “female”, “male” to 0, 1.

```

unique(mydata$previous_loan_defaults_on_file)
mydata$previous_loan_defaults_on_file <- factor(mydata$previous_loan_defaults_on_file,
                                                levels = c("No", "Yes"), labels = c(0,1))
> unique(mydata$previous_loan_defaults_on_file)
[1] "No"    "Yes"

```

```
> unique(mydata$previous_loan_defaults_on_file)
[1] 0 1
Levels: 0 1
```

→ Converted “No”, “Yes” to 0, 1.

## Numeric to Categorical:

```
unique(mydata$loan_status)

mydata$loan_status <- factor(mydata$loan_status,
                             levels = c(0,1),labels = c("No","Yes"))

> unique(mydata$loan_status)
[1] "1" "0"

> unique(mydata$loan_status)
[1] Yes No
Levels: No Yes
```

-> Converted 1,0 to “Yes”, “No”

## Code: Min-Max Normalization

```
normalization <- function(n) {
  (n - min(n, na.rm = TRUE)) / (max(n, na.rm = TRUE) - min(n, na.rm = TRUE))
}

mydata$person_income <- normalization(mydata$person_income)
```

→ This code defines a min-max normalization function to scale numeric values between 0 and 1 and applies it to the person\_income column.

## Code: Filter data

```
filtered_data1 <- mydata %>%
  filter(loan_amnt > 20000 | person_age < 20)

filtered_data1
```

```

str(filtered_data1)
> str(filtered_data1)
'data.frame': 126 obs. of 14 variables:
 $ person_age          : int 21 23 24 22 24 22 23 27 23 ...
 $ person_gender        : Factor w/ 2 levels "1","2": 1 1 2 1 2 1 1 2 2 1 ...
 $ person_education     : chr "Master" "Bachelor" "Master" "Bachelor" ...
 $ person_income         : int 71948 79753 66135 149875 95550 100684 102985 114860 130713 149875 ...
 $ person_emp_exp       : int 0 0 1 5 3 0 3 0 0 ...
 $ person_home_ownership: chr "RENT" "RENT" "RENT" ...
 $ loan_amnt            : int 35000 35000 35000 35000 35000 35000 35000 35000 35000 ...
 $ loan_intent           : chr "PERSONAL" "MEDICAL" "MEDICAL" "EDUCATION" ...
 $ loan_int_rate         : num 16 15.2 14.3 12.4 11.1 ...
 $ loan_percent_income   : num 0.49 0.44 0.53 0.37 0.37 0.35 0.34 0.3 0.27 0.25 ...
 $ cb_person_cred_hist_length: int 3 2 4 3 4 2 4 2 4 4 ...
 $ credit_score          : int 561 675 586 701 585 544 621 573 708 583 ...
 $ previous_loan_defaults_on_file: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 ...
 $ loan_status            : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
> |

```

- This code filters the dataset to include rows where loan\_amnt is greater than 20,000 or person\_age is less than 20, then displays the filtered data and its structure.

```
filtered_data2 <- mydata %>%
```

```
filter(loan_status == "Yes" & person_age > 20)
```

```
filtered_data2
```

```
str(filtered_data2)
```

```

> str(filtered_data2)
'data.frame': 125 obs. of 14 variables:
 $ person_age          : int 21 25 23 24 27 22 24 22 21 22 ...
 $ person_gender        : Factor w/ 2 levels "1","2": 1 1 1 2 1 1 2 1 1 1 ...
 $ person_education     : chr "Master" "High School" "Bachelor" "Master" ...
 $ person_income         : int 71948 12438 79753 66135 12951 149875 95550 100684 12739 102985 ...
 $ person_emp_exp       : int 0 3 0 1 0 1 5 3 0 0 ...
 $ person_home_ownership: chr "RENT" "MORTGAGE" "RENT" "RENT" ...
 $ loan_amnt            : int 35000 5500 35000 35000 2500 35000 35000 35000 1600 35000 ...
 $ loan_intent           : chr "PERSONAL" "MEDICAL" "MEDICAL" "MEDICAL" ...
 $ loan_int_rate         : num 16.02 12.87 15.23 14.27 7.14 ...
 $ loan_percent_income   : num 0.49 0.49 0.44 0.53 0.19 0.37 0.37 0.35 0.13 0.34 ...
 $ cb_person_cred_hist_length: int 3 3 2 4 2 3 4 2 3 4 ...
 $ credit_score          : int 561 635 675 586 532 701 585 544 640 621 ...
 $ previous_loan_defaults_on_file: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 ...
 $ loan_status            : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
>

```

- This code selects rows where loan\_status is "Yes" and person\_age is greater than 20, then shows the filtered dataset and its structure.

## Sampling: Oversampling method

```
table(mydata$person_gender)
```

```
balanced_data_over <- mydata %>%
```

```
group_by(person_gender) %>%
```

```
slice_sample(n = max(table(mydata$person_gender)), replace = TRUE) %>%
ungroup()
```

	person_age	person_gender	person_education	person_income	person_emp_exp	person_home_ownership	loan_amnt	loan_intent	loan_int_rate	loan_precio
1	22	1	Bachelor	79054	1	RENT	32500	HOMEIMPROVEMENT	9.99	
2	22	1	Master	79255	0	RENT	34000	EDUCATION	17.58	
3	24	1	Bachelor	83853	3	RENT	27600	HOMEIMPROVEMENT	11.01	
4	23	1	Associate	79183	0	RENT	28000	EDUCATION	7.90	
5	22	1	Associate	72608	3	RENT	25000	PERSONAL	10.99	
6	23	1	High School	86811	0	RENT	30000	DEBTCONSOLIDATION	11.01	
7	23	1	High School	111369	0	RENT	35000	MEDICAL	20.00	
8	23	1	Associate	117250	0	RENT	30000	VENTURE	10.65	
9	23	1	Bachelor	17149	0	RENT	4950	DEBTCONSOLIDATION	7.90	
10	22	1	High School	68824	0	RENT	25000	HOMEIMPROVEMENT	14.96	

-> This code counts the number of records for each gender and then creates a balanced dataset by oversampling the minority gender to match the count of the majority gender.

```
table(balanced_data_over$person_gender)
class_distribution_balanced <- balanced_data_over %>%
count(person_gender) %>%
mutate(percentage = n / sum(n) * 100)
```

	person_gender	n	percentage
1	1	123	50
2	2	123	50

➔ This code shows the number of records for each gender in the balanced dataset and creates a table with both counts and their percentage of the total.

## Test and Training dataset

```
set.seed(123)
train_data <- mydata %>%
sample_frac(0.8)
test_data <- mydata %>%
anti_join(train_data)
str(train_data)
```

```
str(test_data)
```

-> This code splits the dataset into a training set (80% of the data) and a test set (remaining 20%), then displays the structure of the training set.

```
> str(train_data)
'data.frame': 161 obs. of 14 variables:
 $ person_age      : int 26 22 27 24 25 22 23 24 26 26 ...
 $ person_gender   : Factor w/ 2 levels "1","2": 2 2 2 1 2 1 2 1 1 2 ...
 $ person_education: chr "Associate" "High School" "Master" "Associate" ...
 $ person_income    : int 68066 241030 130713 82741 65742 80838 16599 15082 241114 22674
 $ person_emp_exp  : int 2 3 0 0 3 0 0 0 5 0 ...
 $ person_home_ownership: chr "RENT" "RENT" "RENT" ...
 $ loan_amnt       : int 25000 15000 35000 25000 25000 30000 2400 2500 3000 20000 ...
 $ loan_intent     : chr "DEBTCONSOLIDATION" "EDUCATION" "EDUCATION" "PERSONAL" ...
 $ loan_int_rate   : num 12.9 14.8 18.4 14.1 16.8 ...
 $ loan_percent_income: num 0.37 0.06 0.27 0.3 0.38 0.37 0.14 0.17 0.01 0.09 ...
 $ cb_person_cred_hist_length: int 2 3 4 3 4 3 2 3 3 2 ...
 $ credit_score    : int 637 626 708 688 672 645 505 631 608 621 ...
 $ previous_loan_defaults_on_file: Factor w/ 2 levels "0","1": 1 2 1 1 1 1 1 2 1 ...
 $ loan_status     : Factor w/ 2 levels "No","Yes": 2 1 2 2 2 2 2 2 1 2 ...
```

➔ Train data total 161 obs

```
> str(test_data)
'data.frame': 40 obs. of 14 variables:
 $ person_age      : int 21 23 23 23 27 27 22 26 26 23 ...
 $ person_gender   : Factor w/ 2 levels "1","2": 1 1 2 2 2 2 2 2 2 ...
 $ person_education: chr "High school" "Associate" "High School" "Bachelor" ...
 $ person_income    : int 12739 149875 111369 136628 14293 144855 361547 360977 337133 333399 ...
 $ person_emp_exp  : int 0 0 0 0 0 1 0 5 7 1 ...
 $ person_home_ownership: chr "OWN" "RENT" "RENT" "RENT" ...
 $ loan_amnt       : int 1600 35000 35000 35000 1400 32000 24250 25000 10000 35000 ...
 $ loan_intent     : chr "VENTURE" "EDUCATION" "MEDICAL" "DEBTCONSOLIDATION" ...
 $ loan_int_rate   : num 14.74 7.9 20 18.25 9.32 ...
 $ loan_percent_income: num 0.13 0.25 0.31 0.26 0.1 0.22 0.07 0.07 0.03 0.1 ...
 $ cb_person_cred_hist_length: int 3 4 4 4 3 2 2 3 3 4 ...
 $ credit_score    : int 640 583 694 709 607 586 637 695 623 609 ...
 $ previous_loan_defaults_on_file: Factor w/ 2 levels "0","1": 1 1 1 1 1 2 2 2 2 1 ...
 $ loan_status     : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 1 1 1 1 1 ...
```

➔ Test data total 40 obs

## Descriptive Statistics

```
summarize_age_income<- mydata %>%
  group_by(loan_status) %>%
  summarise(
    mean_age = mean(person_age, na.rm = TRUE),
    median_age = median(person_age, na.rm = TRUE),
    sd_age = sd(person_age, na.rm = TRUE),
    min_age = min(person_age, na.rm = TRUE),
    max_age = max(person_age, na.rm = TRUE),
```

```

mean_income = mean(person_income, na.rm = TRUE),
median_income = median(person_income, na.rm = TRUE),
sd_income = sd(person_income, na.rm = TRUE),
min_income = min(person_income, na.rm = TRUE),
max_income = max(person_income, na.rm = TRUE)

)

```

summarize\_age\_income

```

`summarize_age_income
# A tibble: 2 × 11
  loan_status mean_age median_age sd_age min_age max_age mean_income median_income sd_income min_income max_income
<fct>        <dbl>     <dbl>   <dbl>    <dbl>   <dbl>      <dbl>     <dbl>    <dbl>    <dbl>      <dbl>
1 No          24.1      24       1.76     21       27     232460.    249174   95317.    12282     368115
2 Yes         23.4      23       1.69     21       27     25750.     22608   54691.    12438     316466
|
```

## Average credit\_score between customers with loan\_status 1 and with loan\_status 0

```

avg_credit <- mydata %>%
  group_by(loan_status) %>%
  summarise(
    mean_credit = mean(credit_score, na.rm = TRUE),
    median_credit = median(credit_score, na.rm = TRUE),
    sd_credit = sd(credit_score, na.rm = TRUE),
    count = n()
)

```

avg\_credit

```

> avg_credit
# A tibble: 2 × 5
  loan_status mean_credit median_credit sd_credit count
<fct>           <dbl>        <dbl>      <dbl>    <dbl>
1 No            630.        624.       52.2     76
2 Yes           628.        634        50.0    125
> |
```

## Spread in person\_emp\_exp for customers with different levels of person\_education

```
emp_exp_stats <- mydata %>%
  group_by(person_education) %>%
  summarise(
    mean_exp = mean(person_emp_exp, na.rm = TRUE),
    median_exp = median(person_emp_exp, na.rm = TRUE),
    sd_exp = sd(person_emp_exp, na.rm = TRUE),
    min_exp = min(person_emp_exp, na.rm = TRUE),
    max_exp = max(person_emp_exp, na.rm = TRUE),
    IQR_exp = IQR(person_emp_exp, na.rm = TRUE),
    n = n()
  )

emp_exp_stats
unique(mydata$person_education)

  person_education mean_exp median_exp sd_exp min_exp max_exp IQR_exp n
  <chr>           <dbl>     <dbl>   <dbl>  <int>   <dbl>   <dbl> <int>
1 Associate        1.20      1       1.41    0       5       2       46
2 Bachelor         1.79      1       1.72    0       5       3       73
3 Doctorate        2         2       NA      2       2       0       1
4 High School      1.41      1       1.80    0       7       3       58
5 Master           1.74      1       1.89    0       6       2.5     23
> unique(mydata$person_education)
[1] "Master"      "High School" "Bachelor"    "Associate"   "Doctorate"
```